```
import pandas as pd
import numpy as np
from sklearn.svm import SVC
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler


#Importing the Data from the give files
train_set = pd.read_csv("tweets-train-data.csv", sep= ",", header = None, names = ["Tweet", "Date", "Retweets", "Likes", "Location"])
test_set = pd.read_csv("tweets-test-data.csv", sep= ",", header = None, names = ["Tweet", "Date", "Retweets", "Likes", "Location"])
train_targets = pd.read_csv("tweets-train-targets.csv", sep= ",", header = None, names = ["Author"])
test_targets = pd.read_csv("tweets-test-targets.csv", sep= ",", header = None, names = ["Author"])

#Formating the data
training = pd.concat([train_set, train_targets], axis = 1)
training = training[training["Date"].astype(str).str.startswith("2016")]
training["Likes"] = training["Likes"].astype(float)
training["Date"], training["Time"] = training["Date"].str.split("T").str
training["Date"].astype("datetime64")
print(training)

#Likes information
Trump = training[training["Author"] == "DT"]
Clinton = training[training["Author"] == "HC"]
Clinton_likes = Clinton["Likes"]
Trump_likes = Trump["Likes"]
print(Trump_likes)
print(Clinton_likes)

#Ploting the retweets
Clinton_rt = Clinton["Retweets"]
Trump_rt = Trump["Retweets"]

plt.plot(Clinton_rt, "c")
plt.plot(Trump_rt, "r")
plt.xlabel("Tweet id")
plt.ylabel("Number of rts")
plt.legend(("HC", "DT"))
plt.grid(True)

plt.show()
plt.close()

#Clinton most popular tweet (rts)
print(max(Clinton_rt))
top_Clinton_rt = training.loc[training["Retweets"]==max(Clinton_rt)]
print(top_Clinton_rt)

#Trump most popular tweet (rts)
print(max(Trump_rt))
top_Trump_rt = training.loc[training["Retweets"]==max(Trump_rt)]
print(top_Trump_rt)

#Clinton most liked tweet
print(max(Clinton_likes))
top_Clinton_likes = training.loc[training["Likes"]==max(Clinton_likes)]
print(top_Clinton_likes)

#Trump most liked tweet
print(max(Trump_likes))
top_Trump_likes = training.loc[training["Likes"]==max(Trump_likes)]
print(top_Trump_likes)

#Formating the testing data
testing = pd.concat([test_set, test_targets], axis = 1)
testing = testing[testing["Date"].astype(str).str.startswith("2016")]
testing["Likes"] = testing["Likes"].astype(float)
testing["Date"], testing["Time"] = testing["Date"].str.split("T").str
```

```
69  testing["Date"].astype("datetime64")
70  print(testing)
71
72  #Extracting the features
73  train_features = training["Author"]
74  train_rts = training["Retweets"].values.reshape(-1,1)
75  test_features = testing["Author"]
76  test_rts = testing["Retweets"].values.reshape(-1,1)
77
78
79  #Normalize the tweets length
80  a = np.mean(training["Tweet"].apply(len))
81  b = np.mean(testing["Tweet"].apply(len))
82
83  training["word_length"] = training["Tweet"].apply(len)-a
84  testing["word_length"] = testing["Tweet"].apply(len)-b
85
86  #Re-assign
87  Trump = training[training["Author"] == "DT"]
88  Clinton = training[training["Author"] == "HC"]
89
90  Trump_average_length = Trump["word_length"]
91  Clinton_average_length = Clinton["word_length"]
92
93  #Spaces count
94  training["space_count"] = training["Tweet"].str.count(" ")
95  training["a_count"] = training["Tweet"].str.count("a")
96  testing["space_count"] = testing["Tweet"].str.count(" ")
97  testing["a_count"] = testing["Tweet"].str.count("a")
98
99  #Taking logarithms (monotonic function)
100 training["log_likes"] = np.log(training["Likes"])
101 training["log_rts"] = np.log(training["Retweets"])
102 testing["log_likes"] = np.log(testing["Likes"])
103 testing["log_rts"] = np.log(testing["Retweets"])
104
105 #Dot count
106 training["dot_count"] = training["Tweet"].str.count(".")
107 testing["dot_count"] = testing["Tweet"].str.count(".")
108
109 #Hours
110 training["hour"] = training["Time"].str[0:2]
111 pd.to_numeric(training["hour"])
112 testing["hour"] = testing["Time"].str[0:2]
113 pd.to_numeric(testing["hour"])
114
115 #Months
116 training["month"] = training["Date"].str[5:7]
117 pd.to_numeric(training["month"])
118 testing["month"] = testing["Date"].str[5:7]
119 pd.to_numeric(testing["month"])
120
121 scale = StandardScaler()
122
123 log_training_matrix = training[["log_likes", "log_rts", "word_length", "a_count",
    "space_count", "dot_count", "month", "hour", "Retweets"]]
124 log_testing_matrix = testing[["log_likes", "log_rts", "word_length", "a_count",
    "space_count", "dot_count", "month", "hour", "Retweets"]]
125
126 transformed_x = scale.fit_transform(log_training_matrix, train_features)
127 transformed_y = scale.fit_transform(log_testing_matrix, test_features)
128
129 #Training using SVC
130
131 clf = SVC(C=10, kernel = "rbf", gamma = 0.02)
132 clf.fit(transformed_x, train_features)
133
134 y_predict = clf.predict(transformed_y)
135
136 from sklearn import metrics
137
138 report = metrics.classification_report(test_features, y_predict)
```

```
139 metrics.accuracy_score(test_features, y_predict)
140
141 #K-fold cross Validation
142 try:
143     from sklearn.model_selection import KFold, cross_val_score
144     legacy = False
145 except ImportError:
146     from sklearn.model_selection import KFold, cross_val_score
147     legacy = True
148
149 #K=3
150 if legacy:
151     kf = KFold(len(train_features), n_folds = 3, shuffle = True, random_state = 42)
152 else:
153     kf = KFold(n_splits = 3, shuffle = True, random_state = 42)
154
155 gamma_values = [0.01, 0.001, 0.0001, 0.00001]
156 accuracy_scores = []
157
158 #K-Fold cross validation algorithm:
159 #-Train predictor
160 #-Compute score
161 for gamma in gamma_values:
162
163     #Train classifier
164     clf = SVC(C=10, kernel = "rbf", gamma = gamma)
165
166     #Score
167     if legacy:
168         scores = cross_val_score(clf, transformed_x, train_features, cv = kf,
            scoring="accuracy")
169     else:
170         scores = cross_val_score(clf, transformed_x, train_features, cv =
            kf.split(transformed_x), scoring = "accuracy")
171
172     #Compute Average score
173     accuracy_score = scores.mean()
174     accuracy_scores.append(accuracy_score)
175
176 #Best value of gamma
177 best_index = np.array(accuracy_scores).argmax()
178 best_gamma = gamma_values[best_index]
179
180 #Train with the selected gamma
181 clf = SVC(C=10, kernel = "rbf", gamma = best_gamma)
182 clf.fit(transformed_x, train_features)
183
184 #Evaluate on the test set
185 y_predict = clf.predict(transformed_y)
186 accuracy = metrics.accuracy_score(test_features, y_predict)
187
188 print(accuracy)
189
190 print(best_gamma)
191
192 from sklearn.model_selection import learning_curve
193
194 plt.figure(2)
195 plt.title("Learning curve")
196 plt.xlabel("Training examples")
197 plt.ylabel("Score")
198 plt.grid(True)
199
200
201 clf = SVC(C=10, kernel = "rbf", gamma = best_gamma)
202
203 train_sizes, train_scores, val_scores = learning_curve(clf, transformed_x,
    train_features, scoring = "accuracy")
204
205 train_scores_mean = np.mean(train_scores, axis = 1)
206 train_scores_std = np.std(train_scores, axis = 1)
207 val_scores_mean = np.mean(val_scores, axis = 1)
```

```
val_scores_std = np.std(val_scores, axis = 1)

#Training scores: mean
plt.plot(train_sizes, train_scores_mean, '*' ,color = "r", label = "Training Score")
#Training scores: std
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
train_scores_mean + train_scores_std, alpha = 0.1, color = "r")
#Validation scores: mean
plt.plot(train_sizes, val_scores_mean,'*', color = "c", label = "Cross Validation
Scores")
#Validation scores: std
plt.fill_between(train_sizes, val_scores_mean - val_scores_std, val_scores_mean +
val_scores_std, alpha = 0.1, color = "c")

plt.ylim(0.05, 1.3)
plt.legend()
plt.show()
plt.close()

try:
    from sklearn.model_selection import GridSearchCV
except ImportError:
    from sklearn.grid_search import GridSearchCV
possible_parameters = {
        'C' : [10,4,3,2],
        'gamma' : [0.01,0.001,0.0001,0.00001]
         }

svc = SVC(kernel = "rbf")

#The GridSearch is a classifier. Hence, we try to fit it with the training data,
#and then use it for prediction
clf = GridSearchCV(svc, possible_parameters, n_jobs = 4, cv = 3)
clf.fit(transformed_x, train_features)

y_predidct = clf.predict(transformed_y)
accuracy = metrics.accuracy_score(test_features, y_predict)

#Training phase
clf = SVC(C=10, kernel = "rbf", gamma = 0.01)
clf.fit(transformed_x, train_features)
#Prediction phase
y_predict = clf.predict(transformed_y)

report = metrics.classification_report(test_features, y_predict)
metrics.accuracy_score(test_features, y_predict)


print(report)

#Performance parameters
from sklearn.model_selection import cross_val_score

accuracy = cross_val_score(clf, transformed_x, train_features, cv = 10, scoring =
"accuracy")

#Labels
labels = train_features.map(lambda x: 1 if x == "HC" else 0).values

precision = cross_val_score(clf, transformed_x, labels, cv = 10, scoring =
"precision")

recall = cross_val_score(clf, transformed_x, labels, cv = 10, scoring = "recall")

f1 = cross_val_score(clf, transformed_x, labels, cv=10, scoring = "f1")

print("Avg Precision: {}".format(round(precision.mean(), 3)))
print("Avg Recall: {}".format(round(recall.mean(), 3)))
print("Avg Accuracy: {}".format(round(accuracy.mean(), 3)))
print("Avg f1: {}".format(round(f1.mean(), 3)))
```

```
#Retweets plot
plt.hist(training["Retweets"][training["Author"]=="HC"], color = "c")
plt.hist(training["Retweets"][training["Author"]=="DT"], color = "r")
plt.title("Retweets")
plt.legend()
plt.grid(True)
plt.show()
plt.close()


#Log Retweets plot
plt.hist(training["log_rts"][training["Author"]=="HC"], color = "c")
plt.hist(training["log_rts"][training["Author"]=="DT"], color = "r")
plt.title("Log Retweets")
plt.legend()
plt.grid(True)
plt.show()
plt.close()

#Likes plot
plt.hist(training["Likes"][training["Author"]=="HC"], color = "c")
plt.hist(training["Likes"][training["Author"]=="DT"], color = "r")
plt.title("Likes")
plt.legend()
plt.grid(True)
plt.show()
plt.close()

#Log likes plot
plt.hist(training["log_likes"][training["Author"]=="HC"], color = "c")
plt.hist(training["log_likes"][training["Author"]=="DT"], color = "r")
plt.title("Log Likes")
plt.legend()
plt.grid(True)
plt.show()
plt.close()

```